

```
import pandas as pd
import numpy as np
```

```
df=pd.read_csv('Boston.csv')
print(df)

   Unnamed: 0      crim     zn  indus  chas     nox      rm    age     dis     rad  \
0            1  0.00632  18.0   2.31      0  0.538  6.575  65.2  4.0900      1
1            2  0.02731   0.0   7.07      0  0.469  6.421  78.9  4.9671      2
2            3  0.02729   0.0   7.07      0  0.469  7.185  61.1  4.9671      2
3            4  0.03237   0.0   2.18      0  0.458  6.998  45.8  6.0622      3
4            5  0.06905   0.0   2.18      0  0.458  7.147  54.2  6.0622      3
..          ...
501           502  0.06263   0.0  11.93      0  0.573  6.593  69.1  2.4786      1
502           503  0.04527   0.0  11.93      0  0.573  6.120  76.7  2.2875      1
503           504  0.06076   0.0  11.93      0  0.573  6.976  91.0  2.1675      1
504           505  0.10959   0.0  11.93      0  0.573  6.794  89.3  2.3889      1
505           506  0.04741   0.0  11.93      0  0.573  6.030  80.8  2.5050      1

      tax  ptratio   black  lstat   medv
0    296    15.3  396.90    4.98  24.0
1    242    17.8  396.90    9.14  21.6
2    242    17.8  392.83    4.03  34.7
3    222    18.7  394.63    2.94  33.4
4    222    18.7  396.90    5.33  36.2
..        ...
501   273    21.0  391.99    9.67  22.4
502   273    21.0  396.90    9.08  20.6
503   273    21.0  396.90    5.64  23.9
504   273    21.0  393.45    6.48  22.0
505   273    21.0  396.90    7.88  11.9
```

[506 rows x 15 columns]

```
# Plotting the distribution of important features in the dataset
fig, axs = plt.subplots(2, 3, figsize=(18, 10))

# Feature: 'crim'
sns.histplot(df['crim'], bins=30, kde=True, color='skyblue', ax=axs[0, 0])
axs[0, 0].set_title("Distribution of 'crim' (Crime Rate)")

# Feature: 'zn'
sns.histplot(df['zn'], bins=30, kde=True, color='orange', ax=axs[0, 1])
axs[0, 1].set_title("Distribution of 'zn' (Zoned Land)")


# Feature: 'rm'
sns.histplot(df['rm'], bins=30, kde=True, color='green', ax=axs[0, 2])
axs[0, 2].set_title("Distribution of 'rm' (Rooms per Dwelling)")

# Feature: 'tax'
sns.histplot(df['tax'], bins=30, kde=True, color='red', ax=axs[1, 0])
axs[1, 0].set_title("Distribution of 'tax' (Property Tax Rate)")

# Feature: 'lstat'
sns.histplot(df['lstat'], bins=30, kde=True, color='purple', ax=axs[1, 1])
axs[1, 1].set_title("Distribution of 'lstat' (Lower Status Population)")

# Feature: 'age'
sns.histplot(df['age'], bins=30, kde=True, color='brown', ax=axs[1, 2])
axs[1, 2].set_title("Distribution of 'age' (Building Age)")

plt.tight_layout()
plt.show()
```





```
# Computing the correlation matrix
correlation_matrix = df[numerical_columns].corr()

plt.figure(figsize=(12, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f', linewidths=0.5)
plt.title("Correlation Heatmap of Features", fontsize=16)
```

```
plt.show()
```

Correlation Heatmap of Features



```
missing_values = df.isnull().sum()
missing_values
```

```
0  
Unnamed: 0 0  
crim 0  
zn 0  
indus 0  
chas 0  
nox 0  
rm 0  
age 0  
dis 0  
rad 0  
tax 0  
ptratio 0  
black 0  
lstat 0  
medv 0
```

**dtype:** int64

```
# Drop the 'Unnamed: 0' column  
df = df.drop(columns=['Unnamed: 0'])  
df.head()
```


	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	black	lstat	medv
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98	24.0
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14	21.6
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03	34.7
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5.33	36.2

df.describe()

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	black	lstat	medv
<b>count</b>	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
<b>mean</b>	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	68.574901	3.795043	9.549407	408.237154	18.455534	356.674032	12.653063	22.532806
<b>std</b>	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.148861	2.105710	8.707259	168.537116	2.164946	91.294864	7.141062	9.197104
<b>min</b>	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000	1.129600	1.000000	187.000000	12.600000	0.320000	1.730000	5.000000
<b>25%</b>	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000	2.100175	4.000000	279.000000	17.400000	375.377500	6.950000	17.025000
<b>50%</b>	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.500000	3.207450	5.000000	330.000000	19.050000	391.440000	11.360000	21.200000
<b>75%</b>	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000	5.188425	24.000000	666.000000	20.200000	396.225000	16.955000	25.000000
<b>max</b>	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000	12.126500	24.000000	711.000000	22.000000	396.900000	37.970000	50.000000

```
import seaborn as sns
import matplotlib.pyplot as plt
plt.figure(figsize=(15, 10))
sns.boxplot(data=df)
plt.xticks(rotation=90)
plt.title("Boxplot for Detecting Outliers")
plt.show()
```

Boxplot for Detecting Outliers



```
# Calculate Q1 (25th percentile) and Q3 (75th percentile)
Q1 = df.select_dtypes(include=['float64', 'int64']).quantile(0.25)
Q3 = df.select_dtypes(include=['float64', 'int64']).quantile(0.75)

IQR = Q3 - Q1
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
outliers_iqr = ((df.select_dtypes(include=['float64', 'int64']) < lower_bound) | (df.select_dtypes(include=['float64', 'int64']) > upper_bound))
df_no_outliers_iqr = df[~outliers_iqr.any(axis=1)]
df.shape, df_no_outliers_iqr.shape
```


```
((506, 14), (268, 14))
```

```
# Plot boxplot before removing outliers
plt.figure(figsize=(15, 10))
sns.boxplot(data=df)
plt.xticks(rotation=90)
plt.title("Boxplot Before Outlier Removal")
plt.show()
```

```
# Plot boxplot after removing outliers using IQR method
plt.figure(figsize=(15, 10))
sns.boxplot(data=df_no_outliers_iqr)
plt.xticks(rotation=90)
plt.title("Boxplot After Outlier Removal (IQR)")
plt.show()
```



Boxplot Before Outlier Removal




```
from sklearn.preprocessing import StandardScaler
features_to_scale = df.drop(columns=['chas', 'medv'])

scaler = StandardScaler()
scaled_features = scaler.fit_transform(features_to_scale)
df_scaled = pd.DataFrame(scaled_features, columns=features_to_scale.columns)
df_scaled['chas'] = df['chas']
df_scaled['medv'] = df['medv']
df_scaled.head()
```

	crim	zn	indus	nox	rm	age	dis	rad	tax	ptratio	black	lstat	chas	medv
0	-0.419782	0.284830	-1.287909	-0.144217	0.413672	-0.120013	0.140214	-0.982843	-0.666608	-1.459000	0.441052	-1.075562	0	24.0
1	-0.417339	-0.487722	-0.593381	-0.740262	0.194274	0.367166	0.557160	-0.867883	-0.987329	-0.303094	0.441052	0.492439	0	21.6
2	-0.417342	-0.487722	-0.593381	-0.740262	1.282714	-0.265812	0.557160	-0.867883	-0.987329	-0.303094	0.396427	-1.208727	0	34.7
3	-0.416750	-0.487722	-1.306878	-0.835284	1.016303	-0.809889	1.077737	-0.752922	-1.106115	0.113032	0.416163	-1.361517	0	33.4
4	-0.412482	-0.487722	-1.306878	-0.835284	1.228577	-0.511180	1.077737	-0.752922	-1.106115	0.113032	0.441052	-1.026501	0	36.2

```
# Extract the coefficients (feature importance)
coefficients = pd.DataFrame(model.coef_, X.columns, columns=["Coefficient"])
coefficients["abs_coefficient"] = coefficients["Coefficient"].abs()
coefficients = coefficients.sort_values("abs_coefficient", ascending=False)
plt.figure(figsize=(10, 6))
sns.barplot(x=coefficients.index, y=coefficients['Coefficient'])
plt.title("Feature Importance (Linear Regression)")
plt.xlabel("Feature")
plt.ylabel("Coefficient Value")
plt.xticks(rotation=90)
plt.show()
```





```
from sklearn.model_selection import train_test_split

X = df_scaled.drop(columns=['medv', 'indus', 'age']) # Remove features with low coefficient values
y = df_scaled['medv']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.35, random_state=42)
X_train.shape, X_test.shape, y_train.shape, y_test.shape
```


```
((328, 11), (178, 11), (328, 11), (178, 11))
```

Feature

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
model = LinearRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print(f"Mean Squared Error (MSE): {mse:.2f}")
print(f"R² Score: {r2:.2f}")
```

```
Mean Squared Error (MSE): 20.31
R² Score: 0.73
```

```
plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred, color='blue', edgecolor='k', alpha=0.7)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], color='red', linestyle='--')
plt.title("Actual vs Predicted House Prices")
plt.xlabel("Actual Prices")
plt.ylabel("Predicted Prices")
plt.show()
```



```
independent_vars = ["crim", "zn", "indus", "nox", "rm", "age", "dis", "rad", "tax", "ptratio", "black", "lstat", "chas"]
```

```
results = {}

for col in independent_vars:
    X = df[[col]]
    y = df["medv"]

    model = LinearRegression()
    model.fit(X, y)

    y_pred = model.predict(X)

    # Store actual vs predicted for each feature
    results[col] = pd.DataFrame({
        "Actual_MEDV": y,
        "Predicted_MEDV": y_pred
    })

# Display result keys
list(results.keys())
```

```
['crim',
 'zn',
 'indus',
 'nox',
 'rm',
```

```
'age',
'dis',
'rad',
'tax',
'ptratio',
'black',
'lstat',
'chas']
```

```
results["rm"] # Example: show for RM
```

	Actual_MEDV	Predicted_MEDV
0	24.0	25.175746
1	21.6	23.774021
2	34.7	30.728032
3	33.4	29.025938
4	36.2	30.382152
...	...	...
501	22.4	25.339584
502	20.6	21.034286
503	23.9	28.825691
504	22.0	27.169108
505	11.9	20.215096

506 rows × 2 columns

```
for col in independent_vars:
    print("\n====")
    print(f"Feature: {col}")
    print(results[col])
```

```
Actual_MEDV Predicted_MEDV
0          24.0      29.339845
1          21.6      23.946907
2          34.7      23.946907
3          33.4      22.005449
4          36.2      22.005449
..
501        ...       ...
502        22.4      17.043946
503        20.6      17.043946
503        23.9      17.043946
504        22.0      17.043946
505        11.9      17.043946
```

[506 rows x 2 columns]

```
=====
Feature: black
Actual_MEDV Predicted_MEDV
0          24.0      23.884120
1          21.6      23.884120
2          34.7      23.747396
3          33.4      23.807863
4          36.2      23.884120
..
501        ...       ...
502        22.4      23.719178
503        20.6      23.884120
503        23.9      23.884120
504        22.0      23.768224
505        11.9      23.884120
```

[506 rows x 2 columns]

```
=====
Feature: lstat
Actual_MEDV Predicted_MEDV
0          24.0      29.822595
1          21.6      25.870390
2          34.7      30.725142
```

```
# Cell 2: train/test per feature, metrics, plot & save

summary = []

# Settings
test_size = 0.2
random_state = 42

for col in independent_vars:
    X = df[[col]].values.reshape(-1, 1) # feature as 2D array
    y = df["medv"].values # target

    # train/test split
    X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=test_size, random_state=random_state
    )

    # fit model
    model = LinearRegression()
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)

    # metrics
```

```
r2 = r2_score(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
rmse = np.sqrt(mean_squared_error(y_test, y_pred)) # FIXED RMSE

# Plot Actual vs Predicted (test set)
plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred, color='blue', edgecolor='k', alpha=0.7, s=70, marker='o')

# y = x reference line
mins = min(y_test.min(), y_pred.min()) - 1
maxs = max(y_test.max(), y_pred.max()) + 1
plt.plot([mins, maxs], [mins, maxs], color='red', linestyle='--', linewidth=1.5)


# labels and title
plt.xlabel("Actual Prices (MEDV)", fontsize=12)
plt.ylabel("Predicted Prices (MEDV)", fontsize=12)
plt.title(f"Actual vs Predicted - Feature: {col} | R2={r2:.3f}, MAE={mae:.3f}, RMSE={rmse:.3f}", fontsize=13)
plt.grid(True, linestyle=':', linewidth=0.5)
plt.tight_layout()

# save
filename = f"plots/{col}_actual_vs_pred_test.png"
plt.savefig(filename, dpi=150)
plt.show()


# store summary info
summary.append({
    "feature": col,
    "r2": r2,
    "mae": mae,
    "rmse": rmse,
    "coef": float(model.coef_[0]),
    "intercept": float(model.intercept_),
    "plot_file": filename
})


print(f"Saved plot: {filename} | feature={col} R2={r2:.4f} MAE={mae:.4f} RMSE={rmse:.4f}\n")
```






Saved plot: plots/crim\_actual\_vs\_pred\_test.png | feature=crim R2=0.1587 MAE=5.6311 RMSE=7.8546







Saved plot: plots/zn\_actual\_vs\_pred\_test.png | feature=zn R2=0.2589 MAE=5.3093 RMSE=7.3721




Saved plot: plots/indus\_actual\_vs\_pred\_test.png | feature=indus R2=0.2764 MAE=5.2496 RMSE=7.2845







Saved plot: plots/nox\_actual\_vs\_pred\_test.png | feature=nox R2=0.2052 MAE=5.6954 RMSE=7.6345



Saved plot: plots/rm\_actual\_vs\_pred\_test.png | feature=rm R2=0.3708 MAE=4.4783 RMSE=6.7930





Saved plot: plots/age\_actual\_vs\_pred\_test.png | feature=age R2=0.2400 MAE=5.3357 RMSE=7.4654

Actual vs Predicted — Feature: dis |  $R^2=0.073$ , MAE=5.968, RMSE=8.244

